

# Deep Learning Fundamentals II

Course:  
INFO-6146 Tensorflow & Keras with Python



Developed by:  
Mohammad Noorchenaarboo

May 5, 2025

# Current Section

- 1 History and Key Milestones of Deep Learning
- 2 Advantages Over Traditional Machine Learning
- 3 Neurons, Layers, Activation Functions, and MLPs
- 4 Introduction to ANNs, CNNs, and RNNs

# Why Study the History of Deep Learning?

**Deep learning didn't appear overnight.** It evolved over decades through research breakthroughs, computational advances, and large-scale data availability.

## Purpose

Understanding the history helps contextualize why certain architectures exist and what problems they were designed to solve.

# Early Neural Networks (1950s–1980s)

**Perceptron (1958):** Introduced by Frank Rosenblatt. A simple model for binary classification.

## **Limitations:**

- Could not solve non-linear problems (e.g., XOR)
- Training multi-layer networks was not feasible yet

## **Stagnation**

In the 1970s, interest declined due to limited computational power and algorithmic tools.

# Breakthrough: Backpropagation (1986)

**Backpropagation:** A method for computing gradients in multi-layer networks.

## Key Paper

Rumelhart, Hinton, and Williams (1986): “Learning representations by back-propagating errors”

## Impact

Enabled practical training of multi-layer neural networks, laying the foundation for modern deep learning.

# LeNet (1998): First Convolutional Neural Network

**Yann LeCun et al.** designed LeNet-5 to recognize handwritten digits.

## Key ideas:

- Local receptive fields (convolutions)
- Weight sharing
- Subsampling (pooling)

## Application

Used to process checks by US banks.

# AlexNet (2012): Deep Learning Revival

**AlexNet (Krizhevsky et al.)** won the ImageNet challenge in 2012 with a large CNN.

## Why It Mattered

- Huge performance jump vs. traditional methods
- Trained on GPUs using ReLU, dropout, and data augmentation
- Sparked a wave of deep learning research and applications

# Modern Architectures and Beyond

## Important Milestones:

- **ResNet (2015):** Solved vanishing gradient with skip connections
- **GANs (2014):** Generated realistic images using adversarial training
- **Transformer (2017):** “Attention is All You Need” revolutionized NLP
- **GPT-3 (2020), ViT (2021), Stable Diffusion (2022):** Massive models for generative tasks

## Example

Transformers are now used in text, vision, and audio tasks – including Chat-GPT.



# Loading a Pretrained Model (Hugging Face)

## Using BERT for text classification

```
from transformers import pipeline

classifier = pipeline('text-classification',
                      model='bert-base-uncased')

print(classifier("I love this course!"))
```

## Observation

Modern deep learning leverages pre-trained models to save time and improve accuracy.

# Summary: Key Milestones in Deep Learning

Year & Milestone	Description
1958 – Perceptron	Simple binary classifier; inspired future work but limited
1986 – Backpropagation	Enabled training of multi-layer networks
1998 – LeNet-5	First successful CNN for digit recognition
2012 – AlexNet	ImageNet winner; revived DL with GPU training
2015 – ResNet	Introduced skip connections to solve vanishing gradients
2017 – Transformer	Enabled large-scale parallel NLP with attention
2020+ – GPT, ViT, Diffusion	Scaled DL models to new domains and generative AI

# Current Section

- 1 History and Key Milestones of Deep Learning
- 2 Advantages Over Traditional Machine Learning**
- 3 Neurons, Layers, Activation Functions, and MLPs
- 4 Introduction to ANNs, CNNs, and RNNs

# When Is Deep Learning Better?

**Deep learning** is particularly powerful for complex, high-dimensional, and unstructured data.

## Advantages

- Learns from raw data (e.g., images, text, audio)
- Scales with massive data and compute
- Enables end-to-end modeling without manual features

# Key Advantages of Deep Learning

- 1. Automatic Feature Extraction** Learns hierarchical features directly from data.
- 2. End-to-End Learning** Minimizes need for domain-specific preprocessing.
- 3. Scalability** Performance improves as data and model size grow.

## Informative

Deep learning models improve with data and compute – unlike many traditional algorithms which saturate early.

# Where Traditional ML Still Wins

## Limitations of Deep Learning

- Requires a large dataset
- Slower training and higher compute costs
- Less interpretable

## Strengths of Traditional ML

- Works well on small, tabular datasets
- Models are easier to explain (e.g., decision trees)
- Quick to train and deploy

# Code Comparison: ML vs DL

## Traditional ML with scikit-learn

```
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_features, y)
```

## Deep Learning with Keras

```
from tensorflow.keras import layers, models
model = models.Sequential([
    layers.Dense(128, activation='relu'),
    layers.Dense(1, activation='sigmoid')
])
```

# Comparison: Deep Learning vs. Traditional ML

Aspect	Deep Learning vs. Traditional ML
Feature Engineering	DL learns from raw data; ML requires manual feature crafting
Performance on Unstructured Data	DL excels with text, images, and audio; ML performs poorly without preprocessing
Interpretability	ML is more explainable; DL often acts as a black box
Data Requirement	DL needs large datasets; ML performs well with smaller data
Training Time & Compute	DL needs GPUs/TPUs; ML is lightweight and fast



# Current Section

- 1 History and Key Milestones of Deep Learning
- 2 Advantages Over Traditional Machine Learning
- 3 Neurons, Layers, Activation Functions, and MLPs**
- 4 Introduction to ANNs, CNNs, and RNNs

# What Is an Artificial Neuron?

Inspired by biological neurons, an artificial neuron computes a weighted sum of its inputs and applies an activation function.

$$y = f \left( \sum_{i=1}^n w_i x_i + b \right)$$

## Where:

- $x_i$  = input features
- $w_i$  = weights
- $b$  = bias
- $f$  = activation function (e.g., ReLU, sigmoid)

## Example

If inputs represent exam scores, a neuron could predict pass/fail based on learned weights.

# What Is a Layer?

A **layer** is a group of neurons that process data in parallel. Deep networks stack multiple layers.

## Types of layers:

- **Input layer:** Receives raw data
- **Hidden layers:** Transform data through neurons
- **Output layer:** Produces final prediction

## Informative

Each hidden layer transforms the data space into a more abstract representation.

# Activation Functions: Bringing Non-Linearity

Activation functions help networks learn complex patterns by introducing non-linearity.

## Popular choices:

- ReLU (Rectified Linear Unit)
- Sigmoid
- Tanh
- Softmax (for classification output)

## Example

ReLU allows the network to learn threshold-based patterns:

$$f(x) = \max(0, x)$$

# Python Example: Simple MLP with Keras

## MLP with two hidden layers

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense

model = Sequential([
    Dense(64, activation='relu', input_shape=(20,)),
    Dense(32, activation='relu'),
    Dense(1, activation='sigmoid') # binary classification
])
```

# Example: Sentiment Classification (Text)

## Example

Given a sentence like “This course is amazing!”, an MLP can learn to map text embeddings to a sentiment score (positive/negative).

### Pipeline steps:

- Tokenize and embed text (e.g., using BERT or Word2Vec)
- Feed vector into MLP
- Output a score between 0 and 1 (sentiment)

# Example: Predicting Housing Prices (Tabular Data)

## Example

Train an MLP using features like size, location, number of rooms to predict house price.

**Input:** 10 numerical features **Output:** Continuous price value (regression)

## MLP for regression

```
model = Sequential([
    Dense(128, activation='relu', input_shape=(10,)),
    Dense(1)  # no activation for regression
])
```

# Activation Functions: Summary

Activation Function	Behavior and Use Case
ReLU	Fast and simple. Used in most hidden layers. Output: $\max(0, x)$
Sigmoid	Outputs between 0 and 1. Used for binary classification.
Tanh	Outputs between -1 and 1. Often used in older models.
Softmax	Converts logits into probabilities. Used in multi-class classification.



# Current Section

- 1 History and Key Milestones of Deep Learning
- 2 Advantages Over Traditional Machine Learning
- 3 Neurons, Layers, Activation Functions, and MLPs
- 4 Introduction to ANNs, CNNs, and RNNs**

# What Are Neural Network Architectures?

Different problems require different network structures.

## Three common architectures:

- **ANN (Artificial Neural Network)**: General-purpose, tabular data
- **CNN (Convolutional Neural Network)**: Images and spatial data
- **RNN (Recurrent Neural Network)**: Sequences and time series

## Why This Matters

Choosing the right architecture is key to capturing the structure of your data.

# ANN: Artificial Neural Networks

**ANNs** are feedforward networks that connect every neuron in one layer to every neuron in the next.

**Used for:**

- Tabular data
- Simple regression or classification

## Example

Predicting diabetes using patient health metrics (numerical features).

# CNN: Convolutional Neural Networks

**CNNs** are designed to process grid-like data, such as images.

## Key Components:

- Convolutional layers (extract local features)
- Pooling layers (downsample)
- Fully connected layers (final classification)

## Example

Classify handwritten digits in MNIST dataset using pixel patterns.

## Simple CNN model

```
model = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)),
    MaxPooling2D((2,2)),
    Flatten(),
    Dense(10, activation='softmax')
])
```

# RNN: Recurrent Neural Networks

**RNNs** are suited for sequence data – they maintain internal memory across time steps.

## Used for:

- Time series forecasting
- Natural language processing
- Speech recognition

## Example

Predicting the next word in a sentence or the stock price for tomorrow.

## Simple RNN with Keras

```
model = Sequential([  
    SimpleRNN(64, input_shape=(timesteps, features)),  
    Dense(1)  
)
```

# Comparison: ANN vs CNN vs RNN

Architecture	Best For
ANN	General-purpose structured data (e.g., tabular features)
CNN	Image classification, object detection, spatial pattern recognition
RNN	Time series, sequences, text, speech – where order matters