

# CNN Architecture, OpenCV & Data Augmentation

Course:  
Deep Learning with Tensorflow & Kersa 2



Developed by:  
Mohammad Noorchenaarboo

May 25, 2025

# Current Section

- 1 Convolutional Layers and Kernel Operations
- 2 Stride and Padding Types ('valid' vs. 'same')
- 3 Pooling Layers (Max, Average, Global)
- 4 OpenCV for Image Loading and Preprocessing
- 5 Image Augmentation Using OpenCV and ImageDataGenerator
- 6 Feature Extraction: Canny, Sobel, Harris Corners
- 7 CNN Implementation with TensorFlow + OpenCV Integration

# Why Do We Need Convolution in Neural Networks?

Imagine trying to classify a 100x100 image with a fully connected ANN.

- **Each pixel** is treated as an independent feature.
- This results in  $100 \times 100 = 10,000$  input features – very sparse, non-spatially aware.

## Scalability Warning

Fully connected layers grow *quadratically* with input size – making them unsuitable for large images.

**Problem:** How do we efficiently detect local patterns (edges, textures, etc.) while preserving spatial structure?

**Solution: Convolutional layers**, which use local filters (kernels) and shared weights to learn spatial features.

$$Y(i,j) = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} X(i+m, j+n) \cdot K(m,n)$$

A convolution applies a small kernel (filter) to different regions of the image, performing dot products to extract features.

# Structure of a Convolutional Layer

## Key Parameters:

- **Kernel Size (F):** Size of the filter, e.g.  $3 \times 3$ ,  $5 \times 5$ .
- **Stride (S):** Number of pixels the filter moves each step.
- **Padding (P):** Preserves dimensions using 'same' or reduces them with 'valid'.
- **Depth (D):** Number of filters applied – outputs a volume, not a flat vector.

$$\text{Output Size} = \left\lfloor \frac{W - F + 2P}{S} \right\rfloor + 1$$

## Example: 5x5 Input, 3x3 Filter, No Padding, Stride 1

- Output =  $(5 - 3 + 0)/1 + 1 = 3$
- Resulting feature map =  $3 \times 3$

## Common Pitfall

Incorrect stride or padding can lead to mismatched input/output dimensions in deep CNN stacks.

# Implementing Convolution in Keras

## 2D Convolution Layer in Keras

```
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.models import Sequential

model = Sequential()
model.add(Conv2D(filters=32,
                  kernel_size=(3, 3),
                  strides=(1, 1),
                  padding='same',
                  activation='relu',
                  input_shape=(28, 28, 1)))

model.summary()
# Output shape: (None, 28, 28, 32)
```

## Output Explanation

This layer outputs 32 feature maps of size 28x28, due to 'same' padding.

# Why Shared Weights Matter

**Traditional ANN:** Every neuron has a unique set of weights.

**CNN:** Each kernel is used across the entire image (weight sharing).

$$\text{Total parameters} = (F \times F \times \text{Channels}) \times \text{Number of Filters}$$

## Parameter Efficiency

CNNs drastically reduce the number of trainable parameters, enabling deeper models and faster training.

## Example:

A  $3 \times 3$  kernel on a single-channel image has  $3 \times 3 = 9$  weights, regardless of image size.

## Misconception

Weight sharing does not mean fewer computations – it only reduces parameters, not FLOPs (Floating-point Operations Per Second).

# Summary Table: Convolutional Layers and Kernel Operations

Concept	Explanation
Kernel / Filter	A small weight matrix that moves across the image extracting local features.
Stride	Step size with which the kernel moves – affects output size.
Padding	'Same' keeps output size equal to input; 'Valid' reduces dimensions.
Weight Sharing	Kernel parameters reused across input space – fewer total weights.
Output Volume	For multiple filters, CNN outputs a 3D volume of feature maps.

# Current Section

- 1 Convolutional Layers and Kernel Operations
- 2 Stride and Padding Types ('valid' vs. 'same')**
- 3 Pooling Layers (Max, Average, Global)
- 4 OpenCV for Image Loading and Preprocessing
- 5 Image Augmentation Using OpenCV and ImageDataGenerator
- 6 Feature Extraction: Canny, Sobel, Harris Corners
- 7 CNN Implementation with TensorFlow + OpenCV Integration



# What Happens When Filters Don't Fit Perfectly?

**Question:** How do we handle cases when the filter size does not evenly divide the input image?

$$\text{Output size} = \left\lfloor \frac{W - F + 2P}{S} \right\rfloor + 1$$

## Warning: Spatial Reduction

Without padding, convolving large kernels can rapidly shrink the image – making it unusable in deep CNNs.

**Two padding modes commonly used:**

- **'valid'**: No padding – output shrinks.
- **'same'**: Pads input so that output size equals input size.

## Why Padding Matters

Padding allows deeper models by maintaining spatial resolution, especially in early layers.

# Understanding 'valid' Padding

'valid' padding means no padding at all:

$$\text{Output size} = \left\lfloor \frac{W - F}{S} \right\rfloor + 1$$

**Example: 7x7 Input, 3x3 Filter, Stride 1, 'valid'**

- Output size =  $\left\lfloor \frac{7-3}{1} \right\rfloor + 1 = 5$
- Output feature map =  $5 \times 5$

## Drawback

Each convolution shrinks the feature map, which can lead to loss of border information.

# Understanding 'same' Padding

**'same' padding** ensures that output dimensions match input:

$$P = \left\lfloor \frac{F-1}{2} \right\rfloor \quad (\text{for odd-sized filters})$$

**Goal:** Maintain output size = input size

## Example: 7x7 Input, 3x3 Filter, Stride 1, 'same'

- Padding =  $\left\lfloor \frac{3-1}{2} \right\rfloor = 1$
- Output =  $\left\lfloor \frac{7-3+2 \cdot 1}{1} \right\rfloor + 1 = 7$

## Use Case

Use 'same' padding in early layers to retain spatial resolution for deep stacking.

# Stride: The Step of Convolution

**Stride (S)** defines how far the filter jumps with each move.

- **Stride = 1:** Dense scanning of input.
- **Stride > 1:** Downsamples spatial dimensions.

$$\text{Output size} = \left\lfloor \frac{W - F + 2P}{S} \right\rfloor + 1$$

**Example: 7x7 Input, 3x3 Filter, Padding = 0, Stride = 2**

$$\left\lfloor \frac{7 - 3 + 0}{2} \right\rfloor + 1 = 3 \Rightarrow \text{Output size} = 3 \times 3$$

## Oversized Strides

Large strides may skip important features and lead to aliasing.

# Stride and Padding in Keras

## Experimenting with Stride and Padding

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D

model = Sequential()

# SAME padding
model.add(Conv2D(16, (3,3), strides=(1,1), padding='same',
                input_shape=(28,28,1)))
# VALID padding
model.add(Conv2D(32, (3,3), strides=(2,2), padding='valid'))

model.summary()
# Layer 1 Output: (28, 28, 16)
# Layer 2 Output: (13, 13, 32)
```

## What This Means

'SAME' preserved the input size; 'VALID' with stride 2 reduced spatial dimensions to almost half.

# Summary Table: Padding and Stride

Parameter	Description
'valid' Padding	No padding; output is smaller than input.
'same' Padding	Adds padding so output size equals input size.
Stride = 1	Moves filter by 1 pixel; maximum overlap.
Stride > 1	Skips pixels; reduces spatial size.
Padding and stride interaction	Controls spatial resolution and computational load.

# Current Section

- 1 Convolutional Layers and Kernel Operations
- 2 Stride and Padding Types ('valid' vs. 'same')
- 3 Pooling Layers (Max, Average, Global)**
- 4 OpenCV for Image Loading and Preprocessing
- 5 Image Augmentation Using OpenCV and ImageDataGenerator
- 6 Feature Extraction: Canny, Sobel, Harris Corners
- 7 CNN Implementation with TensorFlow + OpenCV Integration

# Why Do We Pool in CNNs?

**Motivation:** After convolution, feature maps may still be large. Pooling helps:

- Reduce spatial dimensions
- Introduce spatial invariance (translation robustness)
- Control overfitting by downsampling

$$\text{Output size} = \left\lfloor \frac{W - F}{S} \right\rfloor + 1$$

## Definition

**Pooling** is a downsampling operation applied over each feature map independently.

## Types:

- Max Pooling
- Average Pooling
- Global Pooling



# Max Pooling with Numerical Example

**Max Pooling:** Selects the maximum value from each window.

**Input Feature Map (4x4):**

$$\begin{bmatrix} 1 & 3 & 2 & 4 \\ 5 & 6 & 7 & 8 \\ 2 & 1 & 0 & 3 \\ 4 & 9 & 5 & 2 \end{bmatrix}$$

**2x2 Pool, Stride 2  $\Rightarrow$  Output (2x2):**

$$\begin{bmatrix} 6 & 8 \\ 9 & 5 \end{bmatrix}$$

## Interpretation

The highest value in each  $2 \times 2$  block is preserved; other details are discarded.

# Average Pooling with Numerical Example

**Average Pooling:** Computes the average of values in each region.

**Same 4x4 Input:**

$$\begin{bmatrix} 1 & 3 & 2 & 4 \\ 5 & 6 & 7 & 8 \\ 2 & 1 & 0 & 3 \\ 4 & 9 & 5 & 2 \end{bmatrix}$$

**2x2 Pool, Stride 2  $\Rightarrow$  Output (2x2):**

$$\begin{bmatrix} 3.75 & 5.25 \\ 4 & 2.5 \end{bmatrix}$$

## Use Case

Average pooling is often used in tasks requiring smooth spatial reduction (e.g. regression, heatmaps).

## Note

Average pooling may blur out sharp features or edges.

# Global Pooling Layers

## Global Pooling:

- Instead of using a fixed window, apply pooling across the entire feature map.
- Reduces each feature map to a **single value**.

Global Max Pooling:  $y = \max(X)$       Global Average Pooling:  $y = \frac{1}{N} \sum_{i=1}^N x_i$

## When to Use

Used before the final classification layer to flatten feature maps without needing a dense layer.

## Example:

Global max pooling on a  $3 \times 3$  feature map:

$$\begin{bmatrix} 0.5 & 1.2 & 0.3 \\ 0.4 & 2.1 & 0.8 \\ 0.6 & 0.7 & 1.0 \end{bmatrix} \Rightarrow \max = 2.1$$

# Pooling Layers in Keras

## Using Max and Average Pooling in Keras

```
from tensorflow.keras.layers import MaxPooling2D, AveragePooling2D,
    GlobalMaxPooling2D
from tensorflow.keras.models import Sequential

model = Sequential()
model.add(MaxPooling2D(pool_size=(2, 2), strides=2, padding='valid'))
model.add(AveragePooling2D(pool_size=(2, 2), strides=2,
    padding='valid'))
model.add(GlobalMaxPooling2D())

model.summary()
# Outputs:
# MaxPooling2D -> reduces spatial dimensions by 2
# AveragePooling2D -> same reduction
# GlobalMaxPooling2D -> outputs 1 value per feature map
```

## Best Practice

Use max pooling for feature extraction in classification tasks; average pooling for smoother transitions.

# Summary Table: Pooling Layers

Pooling Type	Behavior and Use Case
Max Pooling	Retains the most dominant value in each region; preserves edges.
Average Pooling	Computes the mean of values; useful for smooth generalization.
Global Max Pooling	Reduces each feature map to a scalar by selecting the max value.
Global Average Pooling	Computes average of each feature map; common in modern CNNs.
Stride	Determines how much to shift the pooling window; controls downsampling.

# Current Section

- 1 Convolutional Layers and Kernel Operations
- 2 Stride and Padding Types ('valid' vs. 'same')
- 3 Pooling Layers (Max, Average, Global)
- 4 OpenCV for Image Loading and Preprocessing**
- 5 Image Augmentation Using OpenCV and ImageDataGenerator
- 6 Feature Extraction: Canny, Sobel, Harris Corners
- 7 CNN Implementation with TensorFlow + OpenCV Integration

# Why Use OpenCV in CNN Pipelines?

**Problem:** Deep learning models require standardized image formats and intensities.

**OpenCV (cv2)** is an efficient C++-backed library used for:

- Reading and displaying images
- Converting color spaces
- Resizing, normalizing, and thresholding
- Augmenting training data

## Why OpenCV?

- Fast image I/O and processing
- Consistent preprocessing for training/inference
- Direct NumPy compatibility

CNN Input:  $X \in \mathbb{R}^{H \times W \times C} \Rightarrow$  Standardized via OpenCV

# Reading and Displaying Images

## Loading Image with OpenCV

```
import cv2

img = cv2.imread('image.jpg')           # Loads image in BGR format
cv2.imshow('Original Image', img)
cv2.waitKey(0)
cv2.destroyAllWindows()
```

## Under the Hood

**cv2.imread()** returns a NumPy array of shape  $(H, W, 3)$ :

- 3 channels: Blue, Green, Red
- Data type: unsigned 8-bit integers, range  $[0, 255]$

Pixel Value at  $(x, y)$  = BGR triplet =  $[B, G, R]$

## Note

Images are loaded in BGR format, not RGB – this affects color transformations.



# Grayscale Conversion: What Actually Happens?

`cv2.cvtColor()` can convert a BGR image to grayscale:

## Convert Color to Grayscale

```
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
```

## Mathematics Behind Grayscale

$$\text{Gray} = 0.114 \cdot B + 0.587 \cdot G + 0.299 \cdot R$$

- Based on human eye sensitivity
- Emphasizes green, suppresses blue

# Grayscale Conversion: What Actually Happens?

## Numerical Example

For pixel:  $[B, G, R] = [60, 120, 200]$ :

$$\text{Gray} = 0.114 \cdot 60 + 0.587 \cdot 120 + 0.299 \cdot 200 = 6.84 + 70.44 + 59.8 = 137.08$$

$\Rightarrow$  Rounded = 137

## Caution

This is not a simple average – it is a perceptually weighted transformation.

# Image Resizing: Scale Matters

## Resizing an Image

```
resized = cv2.resize(gray, (64, 64), interpolation=cv2.INTER_AREA)
```

## How It Works

**Resizing** interpolates pixel values to match target dimensions ( $W_{\text{target}}, H_{\text{target}}$ ).

Scaling Ratio =  $\frac{W_{\text{target}}}{W_{\text{orig}}}$     New Pixel = Weighted average of surrounding pixels

## Example:

If resizing a  $128 \times 128$  to  $64 \times 64$ , every  $2 \times 2$  pixel block becomes one.

## Interpolation Tip

Use `cv2.INTER_AREA` for shrinking and `cv2.INTER_CUBIC` or `INTER_LINEAR` for enlarging.

# Pixel Normalization: From 0–255 to 0–1

## Pixel Value Normalization

```
norm_img = resized / 255.0
```

## Why Normalize?

- Neural nets train faster with smaller input ranges.
- Keeps gradients stable during backpropagation.

## Example:

If pixel intensity = 137, then:

$$\text{Normalized} = \frac{137}{255} \approx 0.537$$

## Note

CNNs assume input values in  $[0, 1]$  or standardized (mean = 0, std = 1).

# Complete Pipeline: From BGR to Normalized Tensor

## Image Preprocessing Pipeline

```
import cv2
import numpy as np

img = cv2.imread('image.jpg')           # BGR format
gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY) # Grayscale
resized = cv2.resize(gray, (64, 64))      # Resize
norm = resized / 255.0                   # Normalize

tensor = np.expand_dims(norm, axis=-1)    # Shape: (64, 64, 1)
```

## Final Shape

CNNs expect inputs with shape  $(H, W, C)$  – for grayscale,  $C = 1$ .

# Summary Table: OpenCV Preprocessing Techniques

Technique	Mathematical Interpretation
Grayscale Conversion	Weighted sum: $0.114B + 0.587G + 0.299R$
Resizing	Interpolated mapping from source to target pixel grid
Normalization	$I_{\text{norm}} = \frac{I}{255}$
Tensor Conversion	Ensures shape $(H, W, C)$ by adding a channel axis

# Current Section

- 1 Convolutional Layers and Kernel Operations
- 2 Stride and Padding Types ('valid' vs. 'same')
- 3 Pooling Layers (Max, Average, Global)
- 4 OpenCV for Image Loading and Preprocessing
- 5 Image Augmentation Using OpenCV and ImageDataGenerator**
- 6 Feature Extraction: Canny, Sobel, Harris Corners
- 7 CNN Implementation with TensorFlow + OpenCV Integration

# Why Do We Augment Images in Deep Learning?

## Motivation:

- CNNs can easily overfit on small datasets.
- Augmentation simulates new training samples.
- It improves generalization by exposing the model to more data variation.

## Definition

**Image augmentation** applies transformations (e.g. rotation, scaling) to create altered versions of existing images.

New Image =  $T(I)$  where  $T$  is a transformation matrix

## Warning

Augmentation only applies to training data – not to validation/test sets.



# Geometric Transformations in OpenCV

## Translation:

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix} \cdot \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

### Image Translation with OpenCV

```
import cv2
import numpy as np

tx, ty = 30, 20
M = np.float32([[1, 0, tx], [0, 1, ty]])
translated = cv2.warpAffine(img, M, (img.shape[1], img.shape[0]))
```

### Numerical Example

Pixel at (50, 50) will move to  $(50 + 30, 50 + 20) = (80, 70)$

# Rotation in OpenCV

## Rotation around the center:

$$M = \begin{bmatrix} \cos \theta & -\sin \theta & (1 - \cos \theta)x_c + \sin \theta y_c \\ \sin \theta & \cos \theta & (1 - \cos \theta)y_c - \sin \theta x_c \end{bmatrix}$$

### Rotate an Image

```
center = (img.shape[1]//2, img.shape[0]//2)
M = cv2.getRotationMatrix2D(center, angle=45, scale=1.0)
rotated = cv2.warpAffine(img, M, (img.shape[1], img.shape[0]))
```

### Numerical Example

Rotation by  $45^\circ$  moves point  $(1,0)$  to  $(\cos 45^\circ, \sin 45^\circ) \approx (0.707, 0.707)$

# Flipping, Scaling, and Zooming

## Flipping:

`cv2.flip(img, 1)`  $\Rightarrow$  Horizontal Flip    `cv2.flip(img, 0)`  $\Rightarrow$  Vertical Flip

### Numerical Example: 2x2 Image Array

Original grayscale image:

$$\text{img} = \begin{bmatrix} 100 & 150 \\ 200 & 250 \end{bmatrix}$$

Horizontal flip (`cv2.flip(img, 1)`):

$$\begin{bmatrix} 150 & 100 \\ 250 & 200 \end{bmatrix}$$

Vertical flip (`cv2.flip(img, 0)`):

$$\begin{bmatrix} 200 & 250 \\ 100 & 150 \end{bmatrix}$$

# Flipping, Scaling, and Zooming

## Scaling:

Zoom matrix:  $\begin{bmatrix} s & 0 & 0 \\ 0 & s & 0 \end{bmatrix}$  where  $s > 1$  for zoom-in,  $s < 1$  for zoom-out

### Zooming with Resizing

```
zoomed = cv2.resize(img, None, fx=2, fy=2,
                    interpolation=cv2.INTER_NEAREST)
# Output (approximate visual result):
# [[100 100 150 150]
#  [100 100 150 150]
#  [200 200 250 250]
#  [200 200 250 250]]
```

# Keras: 'ImageDataGenerator' and Random Augmentation

## Keras-Based Augmentation

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator #  
    Import the class for image augmentation  
  
datagen = ImageDataGenerator(                                # Create an instance with  
    the following random transforms:  
    rotation_range=20,                                       # Random rotation in the  
        range [-20°, +20°]  
    width_shift_range=0.2,                                    # Random horizontal shift  
        by up to 20% of image width  
    height_shift_range=0.2,                                  # Random vertical shift by  
        up to 20% of image height  
    zoom_range=0.15,                                         # Random zoom-in or  
        zoom-out by up to 15%  
    horizontal_flip=True,                                    # Random horizontal flip  
        (mirror image)  
    rescale=1./255                                           # Normalize pixel values  
        from [0, 255] to [0, 1]  
)  
  
aug_iter = datagen.flow(X_train, y_train, batch_size=32)  
# Create a generator that yields batches of augmented images and labels
```

# Keras: 'ImageDataGenerator' and Random Augmentation

## How It Works

Each batch is randomly transformed using defined ranges – image dimensions remain consistent.

## Example:

`rotation_range=20` allows rotation in  $[-20^\circ, +20^\circ]$  uniformly at random.

# Summary Table: Augmentation Techniques and Effects

Augmentation	Mathematical Behavior / Purpose
Rotation	Rotation matrix about image center
Translation	Pixel shift: $x' = x + t_x$ , $y' = y + t_y$
Scaling (Zoom)	Multiplies pixel grid by scale factor $s$
Flipping	Flips image along axis (horizontal/vertical)
'ImageDataGenerator'	Randomly applies transformations per batch during training

# Current Section

- 1 Convolutional Layers and Kernel Operations
- 2 Stride and Padding Types ('valid' vs. 'same')
- 3 Pooling Layers (Max, Average, Global)
- 4 OpenCV for Image Loading and Preprocessing
- 5 Image Augmentation Using OpenCV and ImageDataGenerator
- 6 Feature Extraction: Canny, Sobel, Harris Corners**
- 7 CNN Implementation with TensorFlow + OpenCV Integration



# Why Extract Features Before CNNs?

**Motivation:** Before CNNs, edges and gradients were used as engineered features to detect patterns and object boundaries.

## What is a Gradient?

- A gradient measures how much pixel values change across the image.
- Computed in two directions:  $I_x$  (horizontal) and  $I_y$  (vertical).
- It helps detect areas of rapid intensity change – i.e., edges.

$$\nabla I(x, y) = \left[ \frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right] = [I_x, I_y] \Rightarrow \|\nabla I(x, y)\| = \sqrt{I_x^2 + I_y^2}$$

## Why It Matters

The stronger the gradient, the sharper the edge. High gradient magnitude highlights image boundaries, useful in edge detection and feature maps.

# Sobel Edge Detection

**Sobel operator** estimates image gradients in  $x$  and  $y$  directions.

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}, \quad G_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

$$\text{Gradient magnitude: } \sqrt{(G_x * I)^2 + (G_y * I)^2}$$

## Numerical Insight

Gradient is strongest where pixel intensity changes rapidly in horizontal or vertical directions.

# Sobel Edge Detection

## Numerical Example (Manual Calculation)

Input grayscale patch:

$$I = \begin{bmatrix} 10 & 20 & 30 \\ 20 & 40 & 60 \\ 30 & 60 & 90 \end{bmatrix}$$

Sobel filters for edge detection:

$$K_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} \quad K_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Convolve the 3x3 center region:

$$I_x = (-1)(10) + 0(20) + 1(30) + (-2)(20) + 0(40) + 2(60) + (-1)(30) + 0(60) + 1(90) = 160$$

$$I_y = (1)(10) + 2(20) + 1(30) - (1)(30) - 2(60) - 1(90) = -180$$

Compute edge strength:

$$\text{Magnitude} = \sqrt{160^2 + (-180)^2} \approx 240.83$$

# Sobel Edge Detection

## Sobel Filtering in Code (With Detailed Comments)

```
import cv2
import numpy as np

# Define a simple grayscale image manually
img = np.array([[10, 20, 30],
                [20, 40, 60],
                [30, 60, 90]], dtype=np.uint8)  # 8-bit unsigned
                                                # integers

# Compute gradient in X direction using Sobel operator
sobelx = cv2.Sobel(img, cv2.CV_64F, 1, 0, ksize=3)
# cv2.CV_64F means: output will be 64-bit float (to store negative
# values like -180)
# 1 means derivative in x-direction, 0 in y-direction

# Compute gradient in Y direction
sobely = cv2.Sobel(img, cv2.CV_64F, 0, 1, ksize=3)
# 0 in x, 1 in y means vertical gradient

# Combine gradients to compute magnitude (edge strength)
magnitude = np.sqrt(sobelx**2 + sobely**2)

# Central pixel (1,1) expected magnitude: ~240.83
```

# Sobel Edge Detection

## Important Detail: Why Use `cv2.CV_64F`?

The Sobel operator computes derivatives that can be negative. - If the output is stored as `np.uint8` (0 to 255), negative values are clipped to 0. - `cv2.CV_64F` tells OpenCV to use 64-bit floating point numbers, preserving negative values. This is crucial to ensure correct gradient calculation.

## Classical Limitation

Hand-crafted filters like Sobel are static. They cannot learn or adapt like CNN kernels, which are optimized during training.

# Sobel Edge Detection

Term / Parameter	Explanation
$\nabla I(x, y)$	Gradient of intensity: rate of change in pixel values
$I_x, I_y$	Approximated using convolution (e.g., Sobel filters)
<code>cv2.CV_64F</code>	Ensures floating point precision and allows negative values
<code>kernel_size=3</code>	3x3 kernel size used in Sobel filtering
<code>np.sqrt(...)</code>	Combines gradients to compute edge strength at each pixel

# Canny Edge Detection: A Multi-Step Process

## Canny Edge Detection in OpenCV

```
edges = cv2.Canny(gray, threshold1=100, threshold2=200)  
# Detects edges with lower and upper thresholds
```

## Thresholding Logic

- **Strong edges:** gradient magnitude  $> 200 \Rightarrow$  kept
- **Weak edges:** between 100 and 200  $\Rightarrow$  kept only if connected to strong edge
- **Below 100:** suppressed as noise

## Why Use Two Thresholds?

A single threshold would miss weak but important edges. Dual thresholds help retain edge continuity while filtering noise.

# Canny Edge Detection: A Multi-Step Process

## Numerical Example: Hysteresis Thresholding

Gradient magnitude image:

$$\begin{bmatrix} 80 & 120 & 210 \\ 90 & 160 & 180 \\ 50 & 190 & 220 \end{bmatrix} \quad \text{Thresholds: } T_1 = 100, T_2 = 200$$

### Classification:

- Values  $> 200$ : **Strong** edges  $\Rightarrow$  keep (e.g., 210, 220)
- Values  $100 < x < 200$ : **Weak** edges  $\Rightarrow$  maybe keep (e.g., 120, 160, 180, 190)
- Values  $< 100$ : **Discarded** (e.g., 80, 90, 50)

Connected weak pixels (like 190 connected to 220) are preserved.



# Harris Corner Detection

**Corner:** A pixel where intensity changes significantly in both  $x$  and  $y$  directions – unlike edges (one direction) or flat regions (no change).

**Harris Response Function:**

$$R = \det(M) - k \cdot (\text{trace}(M))^2 \quad \text{where } M = \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

$M$  is the second-moment matrix, summarizing how intensity varies in a local patch.

**Small  $k$  (e.g., 0.04)** More sensitive to detecting **corners**; tolerates edges more

**Large  $k$  (e.g., 0.1)** Penalizes edges heavily; fewer false positives, but may miss corners

## Harris Corners in OpenCV

```
gray_float = np.float32(gray)  # Convert to 32-bit float (required)
harris = cv2.cornerHarris(gray_float, blockSize=2, ksize=3, k=0.04)
# blockSize: neighborhood size for corner detection
# ksize: aperture for Sobel (used to compute I_x and I_y)
# k: sensitivity to edges vs. corners (typically 0.04-0.06)
```

# Harris Corner Detection

## Numerical Example: Response Calculation

Let the image patch produce gradients:

$$I_x = 4, \quad I_y = 5 \Rightarrow M = \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} = \begin{bmatrix} 16 & 20 \\ 20 & 25 \end{bmatrix}$$

Now compute:

$$\det(M) = (16)(25) - (20)^2 = 400 - 400 = 0$$

$$\text{trace}(M) = 16 + 25 = 41 \Rightarrow R = 0 - 0.04 \cdot 41^2 = -67.24$$

**Interpretation:** Since  $R < 0$ , this pixel is likely on an edge.

# Harris Corner Detection

## What $R$ Tells Us

- $R \gg 0 \Rightarrow$  strong intensity change in both directions  $\Rightarrow$  **corner**
- $R \approx 0 \Rightarrow$  no change in either direction  $\Rightarrow$  **flat region**
- $R < 0 \Rightarrow$  strong change in one direction only  $\Rightarrow$  **edge**

## Why Use float32?

`cv2.cornerHarris()` expects floating point precision to store non-integer derivatives and intermediate values. Using `uint8` will lead to inaccurate results.

# Visualizing Features and Annotations

## Drawing Features on Image

```
# Mark Canny edges in red
img[edges > 100] = [0, 0, 255]

# Annotate corners
img[harris > 0.01 * harris.max()] = [0, 255, 0]

# Draw a rectangle or circle
cv2.rectangle(img, (10,10), (50,50), (255,0,0), 2)
cv2.circle(img, (60, 60), 10, (0, 255, 255), -1)
```

## Why Visualize?

- Debug preprocessing steps
- Verify model interprets relevant structures
- Annotate features for inspection

# Summary Table: Feature Detection Methods

Method	Mathematical Operation and Use
Sobel	Computes $\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y}$ to estimate edges
Canny	Multi-step pipeline: blur, Sobel, non-max suppression, thresholding
Harris Corners	Measures intensity variation using matrix eigenvalues
Feature Visualization	Overlay results to inspect edge/corner quality and position

# Current Section

- 1 Convolutional Layers and Kernel Operations
- 2 Stride and Padding Types ('valid' vs. 'same')
- 3 Pooling Layers (Max, Average, Global)
- 4 OpenCV for Image Loading and Preprocessing
- 5 Image Augmentation Using OpenCV and ImageDataGenerator
- 6 Feature Extraction: Canny, Sobel, Harris Corners
- 7 CNN Implementation with TensorFlow + OpenCV Integration**

# Complete Workflow: From Raw Image to Prediction

**Objective:** Use OpenCV to preprocess images, then train a CNN on MNIST/CIFAR-10 using TensorFlow.

## Pipeline Steps:

- 1 Load image using OpenCV
- 2 Preprocess: grayscale, resize, normalize
- 3 Reshape for CNN input
- 4 Define and compile CNN model
- 5 Train and evaluate

$X_{\text{CNN}} = \text{OpenCV}(\text{img}) \rightarrow \text{resize} \rightarrow \text{normalize} \rightarrow \text{expand\_dims}$

## Target Shape

CNN expects input shape:  $(N, H, W, C)$

# Preprocessing Pipeline Using OpenCV

## OpenCV Preprocessing Function

```
def preprocess_image(path, target_size=(32, 32)):  
    img = cv2.imread(path)           # Read image from file  
    (in BGR format by default)  
    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)  # Convert image to  
    grayscale  
    resized = cv2.resize(gray, target_size)        # Resize image to target  
    dimensions (default 32x32)  
    norm = resized / 255.0                # Normalize pixel values  
    to range [0, 1]  
    tensor = np.expand_dims(norm, axis=-1)        # Add channel dimension:  
    (H, W) -> (H, W, 1)  
    return tensor                          # Return preprocessed  
    image tensor
```

## Numerical Example

Original: 128x128 RGB → Grayscale → Resize to 32x32 → Normalize → Add channel axis → (32, 32, 1)

## Integration Ready

This function can be used in custom datasets, not just built-in Keras ones.



# MNIST CNN Model with TensorFlow

## CNN for MNIST (Grayscale)

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
    Dense, Dropout
model = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2,2)),
    Conv2D(64, (3,3), activation='relu'),
    MaxPooling2D((2,2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
    loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

## Why It Works

- Two conv-pool blocks
- Dense + dropout layer to control overfitting
- Output: 10 classes (digits 0–9)

# CIFAR-10 CNN Model (RGB)

## CNN for CIFAR-10 (Color)

```
model = Sequential([
    Conv2D(32, (3,3), activation='relu', input_shape=(32, 32, 3)),
    Conv2D(32, (3,3), activation='relu'),
    MaxPooling2D(pool_size=(2,2)),
    Dropout(0.25),

    Conv2D(64, (3,3), activation='relu'),
    MaxPooling2D(pool_size=(2,2)),
    Dropout(0.25),

    Flatten(),
    Dense(512, activation='relu'),
    Dropout(0.5),
    Dense(10, activation='softmax')
])
```

## Design Insight

Deeper than MNIST model due to complexity of color images and 10 general object classes.

# Train, Evaluate, and Integrate

## Compile and Train Model

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])  
model.fit(X_train, y_train, batch_size=64, epochs=10,  
          validation_split=0.2)
```

## Model Evaluation

Use `model.evaluate()` to test on unseen data. You can also use OpenCV-processed inputs.

## OpenCV + CNN Integration

You can feed any OpenCV-processed image directly into a trained CNN:

```
x = preprocess_image('custom_image.jpg', (28,28))  
x = np.expand_dims(x, axis=0) # Add batch dim  
prediction = model.predict(x)
```

# Summary Table: TensorFlow + OpenCV CNN Pipeline

Step	Technique and Explanation
Load & Convert Image	OpenCV reads image in BGR; convert to grayscale or RGB
Resize & Normalize	Resize to fixed shape (28x28 or 32x32); normalize to [0,1]
CNN Model (MNIST)	2 Conv layers + 1 Dense; input shape (28,28,1)
CNN Model (CIFAR-10)	More conv/pool layers; input shape (32,32,3)
Evaluation	Predict class from OpenCV-preprocessed image