

# Gradient Descent and Its Variants

Course:  
INFO-6154 Data Science and Machine Learning



Developed by:  
Mohammad Noorchenarboo

May 3, 2025

- 1 Gradient Descent and Its Variants
  - Why We Need Gradient Descent
  - What Is a Gradient?
  - Basic Gradient Descent Algorithm
  - Types of Gradient Descent
  - The Role of the Learning Rate
  - Momentum
  - Practical Guidelines and Summary

# Why Do We Need Gradient Descent?

In the previous session, you optimized a model by manually guessing parameters.

But in real applications, this becomes impossible:

- 1,000+ features instead of 1
- Millions of data points
- Non-linear models with millions of parameters

## Limitation

Manual tuning doesn't scale. We need a **mathematical and automated** solution to improve the model iteratively.

## Real-World Example: Insurance Premium Prediction

- Predict costs based on features: age, BMI, smoking, region, etc.
- Model: linear regression or shallow neural network.
- Goal: minimize error across thousands of historical records.

# Why Do We Need Gradient Descent?

## What Do We Need?

An algorithm that can:

- Evaluate how far off the predictions are
- Compute how to change the parameters to improve
- Repeat this process efficiently over data

## This is the Role of Gradient Descent

Gradient descent uses the slope (gradient) of the loss function to adjust model parameters and minimize error.

# Why Do We Need Gradient Descent?

## How Optimization Appears in Practice

```
# PyTorch
optimizer = torch.optim.SGD(model.parameters(), lr=0.01)
loss.backward()
optimizer.step()

# TensorFlow (Keras)
model.compile(optimizer='sgd', loss='mse')
model.fit(X_train, y_train, epochs=10)
```

# What Is a Gradient? (Conceptual View)

The **gradient** is a mathematical tool that tells us which direction to move in order to reduce a function's value the fastest.

**In 1D:** The gradient is the slope of the function.

$$\text{If } L(w) = (2w - 4)^2, \text{ then } \frac{dL}{dw} = 4(2w - 4)$$

- If the gradient is **positive**, the slope is going up – move **left**.
- If the gradient is **negative**, the slope is going down – move **right**.


## Real-World Analogy:

- Imagine you're on a hill with fog and no map.
- The gradient tells you the direction of steepest uphill.
- To go downhill and reach the minimum, you move in the opposite direction.

# What Is a Gradient? (Conceptual View)

## Takeaway

In optimization, we use the **negative gradient** to update model parameters and reduce the loss.



# What Is a Gradient? (In Code)

Let's see how gradients are computed automatically using PyTorch's auto-differentiation.

## Computing Gradients in PyTorch

```
import torch

w = torch.tensor(2.0, requires_grad=True)
loss = (2 * w - 4)**2

loss.backward()
print(w.grad)  # Output: 8.0
```



# What Is a Gradient? (In Code)

## Explanation:

- The function is:  $(2 * w - 4)^2$
- When  $w = 2$ , then  $(2 * 2 - 4)^2 = 0$
- The gradient is:  $dL/dw = 4(2w - 4)$ , so when  $w = 2$ , it's 8

The gradient tells us how the loss would change if we slightly change the parameter  $w$ . Since the gradient is positive, decreasing  $w$  would reduce the loss—this is the core idea behind gradient descent.

## Auto-Differentiation

Modern libraries like PyTorch and TensorFlow compute gradients automatically using a technique called **auto-diff**.

# Gradient Descent: The Core Idea

Once we know the gradient, we can update the model's parameters to reduce the loss.

## Gradient Descent Update Rule

$$\theta \leftarrow \theta - \eta \cdot \nabla_{\theta} L(\theta)$$

### What each term means:

- $\theta$  – model parameters (e.g., weights and biases)
- $\eta$  – learning rate (step size)
- $\nabla_{\theta} L(\theta)$  – gradient of the loss with respect to the parameters

**Key idea:** Move in the direction that decreases the loss.

## Reminder

We use the **negative gradient** because we want to minimize the loss.

# Numerical Example: One Step of Gradient Descent

Suppose we have:

$$L(w) = (2w - 4)^2$$

Then:

$$\frac{dL}{dw} = 4(2w - 4)$$

Assume:

- Initial weight:  $w = 1$
- Learning rate:  $\eta = 0.1$

## Step-by-step:

- Compute gradient:  $4(2 \cdot 1 - 4) = -8$
- Update:  $w \leftarrow 1 - 0.1 \cdot (-8) = 1 + 0.8 = 1.8$

## After One Step

The parameter moved from 1 to 1.8 – we are heading toward the minimum.

# Real-World Analogy

Imagine pushing a shopping cart down a slope.

- The slope tells you which direction the cart will roll.
- A steeper slope means faster movement (larger gradient).
- You adjust the force (step size) with your foot – that's the learning rate.

## Link to Machine Learning

- Cart = model parameters
- Slope = gradient
- Speed of adjustment = learning rate

# Manual Gradient Descent in Python

Let's simulate gradient descent manually for a few steps.

## Manual Gradient Descent (1D)

```
w = 1.0                # Initial value
lr = 0.1               # Learning rate

for step in range(3):
    grad = 4 * (2 * w - 4)
    w = w - lr * grad
    print(f"Step {step + 1}: w = {w}")
```


### Expected Output:

- Step 1:  $w = 1.8$
- Step 2:  $w = 2.28$
- Step 3:  $w = 2.568$

# Manual Gradient Descent in Python

## Observation

Each step moves the parameter closer to the value that minimizes the loss.



# Types of Gradient Descent

Gradient Descent can be applied in different ways depending on how much data is used in each update step.

## The Three Main Types:

- **Batch Gradient Descent** – Uses the entire dataset per update
- **Stochastic Gradient Descent (SGD)** – Uses a single sample per update
- **Mini-Batch Gradient Descent** – Uses a small batch of samples (e.g., 32–256)

# Types of Gradient Descent

## What Changes?

All three follow the same update rule:

$$\theta \leftarrow \theta - \eta \cdot \nabla_{\theta} L(\theta)$$

The only difference is how the gradient is estimated: from the full data, one sample, or a small batch.



# Batch Gradient Descent

## Description:

- Computes the gradient using the entire training set.
- Very stable and smooth updates.
- Computationally expensive for large datasets.

## Example Use Case:

- Small tabular datasets (e.g., diabetes dataset with 300 samples).

### Pros

Very stable convergence – ideal for theoretical analysis.

### Cons

Slow and memory-intensive for large datasets.

# Stochastic Gradient Descent (SGD)

## Description:

- Uses only one training example per update.
- Very noisy, but can escape local minima.
- Often needs careful learning rate tuning and regularization.

## Example Use Case:

- Online learning or streaming data (e.g., real-time fraud detection).

### Pros

Fast updates, good for large-scale or streaming data.

### Cons

High variance in updates – unstable without proper tuning.

# Mini-Batch Gradient Descent

## Description:

- Uses a small batch of data (commonly 32, 64, or 128 samples).
- Combines benefits of both batch and stochastic approaches.
- Most widely used variant in practice.

## Example Use Case:

- Deep learning training on GPUs (e.g., image classification with CIFAR-10).

### Pros

Faster than batch; more stable than SGD; highly parallelizable.

### Cons

Still depends on batch size and tuning learning rate.

# Comparison Table

Method	Data Used per Step	Speed	Stability
Batch GD	All data	Slow	Very stable
SGD	One sample	Fast	Very noisy
Mini-Batch GD	Small batch	Medium	Medium-high

## In Practice

Mini-Batch Gradient Descent is the standard approach in modern deep learning.

## Note

Mini-batch training uses a `DataLoader` to iterate through the dataset in chunks.

# Example: Batch vs. Mini-Batch Code

## Training Loop Comparison

```
# Batch Gradient Descent
for epoch in range(epochs):
    y_pred = model(X_train)
    loss = loss_fn(y_pred, y_train)
    loss.backward()
    optimizer.step()

# Mini-Batch Gradient Descent
for epoch in range(epochs):
    for batch_X, batch_y in dataloader:
        y_pred = model(batch_X)
        loss = loss_fn(y_pred, batch_y)
        loss.backward()
        optimizer.step()
```

# What Is the Learning Rate?

The **learning rate** ( $\eta$ ) controls how big a step we take in the direction of the negative gradient.

$$\theta \leftarrow \theta - \eta \cdot \nabla_{\theta} L(\theta)$$

## Why it matters:

- If  $\eta$  is too **small** – learning is very slow.
- If  $\eta$  is too **large** – updates may overshoot or diverge.
- If  $\eta$  is just right – smooth convergence to minimum.

## Key Insight

The learning rate is one of the most important hyperparameters in optimization.

# Visualizing the Effect of Learning Rate

## Imagine a bowl-shaped loss surface:

- A small  $\eta$  means taking tiny steps downhill.
- A large  $\eta$  might jump across the bowl and oscillate.
- A well-chosen  $\eta$  follows a smooth curved path to the minimum.

## Real-world analogy:

- You're trying to descend a mountain trail in fog.
- If your steps are too short, you'll take forever to get down.
- If they're too long, you might fall or miss the trail.

### What You Want

A step size that adapts to the slope and reaches the goal efficiently.

# Numerical Experiment: Different Learning Rates

Try three different learning rates to see how they behave.

## Simulating Learning Rate Behavior

```
w = 1.0
for step in range(10):
    grad = 4 * (2 * w - 4)
    w = w - 0.01 * grad # Try 0.01, 0.1, 1.0
    print(f"Step {step+1}: w = {w}")
```

## Observation:


- $\eta = 0.01$ : very slow movement
- $\eta = 0.1$ : fast and smooth convergence
- $\eta = 1.0$ : diverges or oscillates wildly



# Numerical Experiment: Different Learning Rates

## Warning

Poor learning rate selection can waste computation or ruin optimization.



# Why Use Momentum?

Even with the right learning rate, gradient descent can struggle in narrow or curved regions of the loss surface.

## Problem: Zig-Zagging

- In steep but narrow valleys, gradients keep changing direction.
- Updates may oscillate back and forth instead of moving forward.


## Solution: Momentum

Instead of updating based only on the current gradient, we **accumulate past gradients** to build velocity.


## Effect:

- Speeds up movement along long slopes
- Dampens oscillations in noisy or curved regions

# Why Use Momentum?



Adapted from [1] (CC-BY 2.0).



Adapted from [1] (CC-BY 2.0).

# How Momentum Works

Momentum introduces a velocity term  $v_t$  that accumulates a moving average of gradients.

$$v_t = \beta v_{t-1} + \eta \cdot \nabla_{\theta} L(\theta)$$

$$\theta \leftarrow \theta - v_t$$

## Where:

- $\beta$  is the momentum coefficient (e.g., 0.9)
- $v_t$  acts like inertia, carrying updates forward

## When to Use It

Momentum helps when:

- You see slow progress in flat regions
- You observe zig-zag behavior in loss curves

# Real-World Analogy: Rolling Ball

## Imagine:

- You're rolling a ball down a hill.
- On flat terrain, it moves slowly.
- On a steep slope, it picks up speed.
- If the path curves, it doesn't change direction immediately – it has momentum.

## Optimization Analogy

- The ball = parameter values
- The hill = loss surface
- Gravity + inertia = gradient + momentum

**Benefit:** Momentum accelerates learning in consistent directions and reduces noise in updates.

# Gradient Descent With and Without Momentum

## Manual Momentum Simulation

```
w = 1.0, v = 0.0, eta = 0.1, beta = 0.9
for step in range(5):
    grad = 4 * (2 * w - 4)
    v = beta * v + eta * grad
    w = w - v
    print(f"Step {step+1}: w = {w}")
```

## Observation:

- Early steps may be small, but velocity builds up.
- Fewer oscillations than regular gradient descent.

## Tuning Tip

Start with  $\beta = 0.9$  and  $\eta = 0.01$  to  $0.1$  depending on the task.

# When to Use Which Gradient Descent Variant?

- **Batch Gradient Descent** – Use for very small datasets where memory is not a concern.
- **SGD (Stochastic)** – Use for online learning or streaming data.
- **Mini-Batch Gradient Descent** – Use as the **default choice** in almost all deep learning tasks.

## Default Starting Points (Recommended)

- Learning rate  $\eta = 0.01$  to  $0.1$
- Mini-batch size = 32 to 128
- Momentum  $\beta = 0.9$

# Common Mistakes to Avoid

## Watch Out For

- Using too high a learning rate without testing convergence
- Not shuffling data during mini-batch training
- Confusing gradient noise with model instability
- Forgetting to zero out gradients between steps (in PyTorch:  
`optimizer.zero_grad()`)

## Good Habits

- Always visualize your loss curve
- Try learning rate scheduling if stuck
- Start simple, then experiment with optimizers



# Session Summary: Gradient Descent and Variants

## In this session, you learned:

- Why manual tuning doesn't scale to real-world ML
- What a gradient is, and how it points toward error reduction
- How basic gradient descent works mathematically and numerically
- The difference between batch, stochastic, and mini-batch updates
- How learning rate and momentum impact convergence

## Next Steps

In the next session, we will explore popular **adaptive optimizers** like Adam, RMSProp, and learning rate schedulers that help training deep models converge faster and more reliably.

# References

- [1] Gabriel Goh.  
Why momentum really works.  
*Distill*, 2017.